VERDIKT · 2026

Release intelligence *for AI teams that ship fast.*

Deterministic certification + confidence-backed recommendation + post-deploy alignment. Verdikt turns release decisions into a closed-loop system.

AI PRODUCT TEAM

HEAD OF AI / ML

CTO · FOUNDER

THE QUESTION THAT CHANGES THE CONVERSATION

> **"If a release degraded in production tomorrow, could you show what risk was flagged before shipping, what the quality scores were, who signed off, and why?"**

If the answer is no, that is the pitch. Not Verdikt. That question.

THE GAP

The eval tooling exists. *The intelligence loop doesn't.*

AI product teams already run Braintrust, LangSmith, hallucination checks, and latency monitoring. **The signals exist.** What is missing is the layer that decides, explains confidence, and validates prediction versus production reality.

What you have

Eval scores

Braintrust, LangSmith, and APM tell you the numbers. None of them issue a verdict or make anyone accountable for the decision to ship.

What's missing

A release intelligence layer

No confidence-backed recommendation, no alignment tracking, and no calibrated learning from misses or over-blocks.

The consequence

Decision quality gap

Overrides are informal, predictions are not validated post-deploy, and the system cannot improve confidence over time. That is the gap Verdikt closes.

THE COST OF NOT HAVING IT

This is already *in your AI incident history.*

A model update ships without an eval gate. Quality drops. Nobody sees it until users complain — three releases later.	SILENT AI REGRESSION — weeks of user-facing degradation before detection, no record of who approved it
An AI feature degrades in production. Engineers reconstruct eval scores, model version, and who approved the release from Slack and eval logs.	HOURS OF SENIOR ENG TIME — per incident, every incident, permanently avoidable with a 30-second lookup
An eval threshold fails. The release is time-sensitive. Someone says "let's ship." No name on the decision. No written justification. No record it happened.	INFORMAL AI OVERRIDE — risk accepted without accountability, nothing documented
An enterprise customer or investor asks what your release quality governance looks like. The honest answer is: informal eval runs and gut feel.	NO DEFENSIBLE RECORD — a critical blocker in serious procurement and Series B+ fundraising conversations

THE SOLUTION

A closed-loop system for every *AI release decision.*

01 — Decide

Verdict against your thresholds

You set the quality standard. Verdikt evaluates every AI signal against it and issues an unambiguous verdict.

02 — Explain

Confidence + reasoning

Verdikt returns a confidence score, risk band, and structured reasoning, with action recommendations for rollout decisions.

03 — Validate

Production alignment

Post-deploy outcomes are inferred from VCS activity (reverts, hotfixes, incident PRs) and compared to pre-release prediction.

04 — Learn

Calibration over time

MISS and OVER_BLOCK patterns adjust future confidence and drive threshold suggestions. The system gets more accurate release by release.

CERTIFIED All signals met threshold. Release is cleared to ship.

UNCERTIFIED One or more signals below threshold. Override required to proceed.

CERTIFIED WITH OVERRIDE Shipped below threshold. Named owner. Written justification. On record.

OVERRIDE ACCOUNTABILITY

The approver can be human or policy. *The record is always there.*

- **Named owner.** Every below-threshold ship requires a named individual with title. Risk is owned by a person, not absorbed anonymously by the team.
- **Mandatory written justification.** Scope of impact. Committed mitigations. Documented rationale — entered at the moment of decision, not reconstructed later.
- **Configurable override authority.** You decide who can override — Tech Lead, Release Manager, QE Lead, VP Engineering, CTO. Configure approval authority to match your org. QE defines the standard; they are not the bottleneck.
- **Permanent, non-editable audit record.** What was written at the moment of override remains on record. Changes are append-only and tamper-evident.
- **Release-ready accountability model.** Model, prompt, and configuration changes move quickly, where ownership can blur. Verdikt records the approver as a named person or role with explicit justification.

Once an organisation requires named override accountability, they do not remove it. Removing it means publicly choosing to remove accountability.

A RELEASE WITH VERDIKT

The same release process.
With intelligence that closes the loop.

01 CI completes

Your existing pipeline runs as normal. Tests, builds, and checks continue unchanged.

02 Signals aggregated

Verdikt pulls quality data from connected sources — Braintrust, LangSmith, Datadog, and Sentry. Eval signals are treated as first-class release inputs.

03 Thresholds evaluated

Each signal is evaluated against your defined thresholds across four categories.

04 Verdict + recommendation issued

Deterministic status plus confidence score, structured reasoning, and recommended rollout action. Unambiguous, attributable, and timestamped.

05 Production alignment computed

After deploy, Verdikt infers production health from VCS signals and classifies CORRECT, MISS, or OVER_BLOCK. Confidence is recalibrated for future releases.

No process disruption. No tooling replacement. This is not another dashboard; it is a release intelligence system attached to your existing flow.

POSITIONING

Sits above your stack.
Replaces nothing. Connects everything.

VERDIKT DOES NOT

- × Replace CI or your build pipeline
- × Replace observability (Datadog, Sentry)
- × Replace test tooling (BrowserStack, Braintrust)
- × Replace release management or deployment
- × Add new testing or eval steps

VERDIKT DOES

- ✓ Issue deterministic certification for each release from your existing signals
- ✓ Generate confidence-backed recommendations with structured reasoning and actions
- ✓ Infer post-deploy outcomes via VCS and compute CORRECT / MISS / OVER_BLOCK
- ✓ Track loop-readiness and tune thresholds with historical simulation and learning signals

GETTING STARTED

Setup in 30 minutes.
No process disruption.

Step 1

Define thresholds

Your QE Lead sets quality baselines per signal category. Defaults provided — adjust to match your current standards.

Step 2

Connect signals + idempotency

Integrate existing tools and CI ingest with idempotency keys so retries never corrupt verdict state.

Step 3

Shadow mode + simulation

Run parallel certification and simulate threshold changes against release history before enforcing gates.

Step 4

Certification live + loop monitor

Gate enabled, recommendations active, VCS monitoring windows open, and loop-readiness tracked in the intelligence hub.

Who configures it: **AI Product Lead or Head of QE** owns threshold setup. Minimal engineering lift for onboarding. No release pipeline redesign required.

WHY NOW

Four pressures that are *already in your roadmap.*

◐ **Release velocity is accelerating**

Model updates and prompt changes now ship multiple times a week. Without a formal decision layer, quality regressions reach users before anyone notices — and nobody clearly owns the risk.

⊞ **Quality fails silently**

A model regression rarely throws an exception or fails CI. It degrades user experience quietly across sessions until someone notices later.

◈ **AI accountability is increasingly ambiguous**

In AI teams, "who approved this change" becomes murky as delivery speeds up. Traditional named-human approval often breaks down. Accountability is the unresolved gap.

⊹ **Enterprise AI scrutiny is increasing**

Enterprise customers and Series B+ investors now ask how you govern release quality. "We have eval runs and code review" is not a defensible answer. Governance evidence is increasingly expected in security, diligence, and procurement conversations.

WHO IT'S FOR

Built for the people *accountable for release quality.*

QE LEADER — CHAMPION / OPERATIONAL OWNER

QE Leader

AI features ship fast and often. When quality degrades, your team is pulled into incidents, while the decision record is scattered across tools and chat.

Defines release standards and enforces one certifiable process. Every below-threshold ship has named ownership on record, so silent regressions stop being invisible.

VP ENGINEERING — ECONOMIC BUYER

VP Engineering

Accountable for what ships, but the quality bar often lives in someone's head, not in a system. When a model update breaks production, reconstruction is slow and brittle.

Gets a documented quality record for every release. Incident reconstruction starts with facts: eval scores, threshold state, who accepted risk, and why.

CTO · FOUNDER — STRATEGIC CHAMPION

CTO / Founder

Enterprise customers ask about AI governance. Investors ask how you certify model quality. Informal eval runs and gut feel are no longer enough.

Every release is certified against defined thresholds. Every override is accountable and on record. A defensible quality process becomes a commercial advantage.

Entry point: One AI product. One release cycle. The AI Product Lead or Head of QE owns setup. Setup takes 30 minutes.

CONFIDENCE COMPOUNDING

Engineering decides.
The whole organisation gains confidence.

> **"When releases are unstable, customer-facing teams absorb the impact first. By the time we hear about it, trust is already damaged."**

WHERE RELEASE RISK SHOWS UP

- **Customer Success** — explains instability to clients
- **Implementation** — manages fallout during onboarding
- **Support** — handles user reports before eng knows
- **Sales** — defends product quality in deals

They are not release decision-makers, but they feel bad releases first.

HOW CONFIDENCE COMPOUNDS

A defensible answer

"This release was certified against agreed quality standards before it shipped."

A documented record

Who approved it, what the quality state was, and what the justification was — immediately accessible.

Institutional confidence

Trust in releases stops being blind faith. It becomes evidence-backed.

THE ASK

One project.
One release cycle. One closed loop.

Run Verdikt on one project for one release cycle. At the end, you will have deterministic certification, confidence-backed recommendations, production alignment, and a measurable loop-readiness baseline.

WHAT WE NEED	Access to your QE Lead and one project to run alongside your existing release process.
WHAT YOU GET	A live release intelligence loop: certification, recommendation confidence, production alignment, and calibration data for better future decisions.
WHAT DOESN'T CHANGE	Your pipeline. Your tooling. Your release cadence. Nothing is replaced. Verdikt sits above it.

verdikt.io · hello@verdikt.io